

AI & ML Agent for Predicting and Remediating Kubernetes Cluster Issues

Presented by ***Capital W***

Problem Statement

Kubernetes clusters often face issues like pod failures, resource bottlenecks, and network disruptions, which impact performance and reliability. Manually identifying and fixing these problems is time-consuming and inefficient. Our solution aims to leverage AI/ML to predict such issues in advance and automate remediation actions, enabling proactive, self-healing Kubernetes environments and improving overall system resilience.

02

Project vision and mission

Our goal is to build an AI/ML-powered system that predicts these failures in advance and triggers automated remediation to ensure reliable and self-healing Kubernetes environments.

01.

Phase 1:

AI/ML model to analyze K8s metrics and predict potential failures.

02.

Phase 2:

Automated or recommended actions to mitigate predicted issues.

Tools Used:

Python, Prometheus,
Kubernetes APIs, Scikit-learn

Our Methodology

We developed a prediction model that leverages both historical and simulated Kubernetes metrics such as CPU usage, memory consumption, pod status, and network I/O. By applying machine learning techniques like time-series analysis and anomaly detection, our model is able to proactively detect potential system failures before they occur. The model specifically focuses on predicting issues such as pod or node crashes, resource exhaustion (CPU, memory, and disk), and network disruptions, which are common and critical in Kubernetes environments.

To complement the prediction system, we designed a multi agent system capable of intelligently predicts and remediates itself. The agent can automatically scale pods when resource usage crosses thresholds, restart or relocate faulty pods to ensure application continuity, and optimize CPU or memory allocation dynamically to prevent performance degradation. This two-fold system not only identifies problems in advance but also ensures rapid and effective resolution, enabling a more reliable and self-healing Kubernetes infrastructure.



04

■ Predictive Failure Modeling

- LSTM models detect pod/node crashes
- Anomaly detection for resource utilization patterns
- 93-94% accurate failure prediction
- Early warning system for imminent failures
-

■ Smart Resource Optimization

- Forecasted usage trends optimize allocation
- KEDA-powered scaling based on custom metrics
- Reduced overprovisioning by 30%
- Workload-aware resource distribution

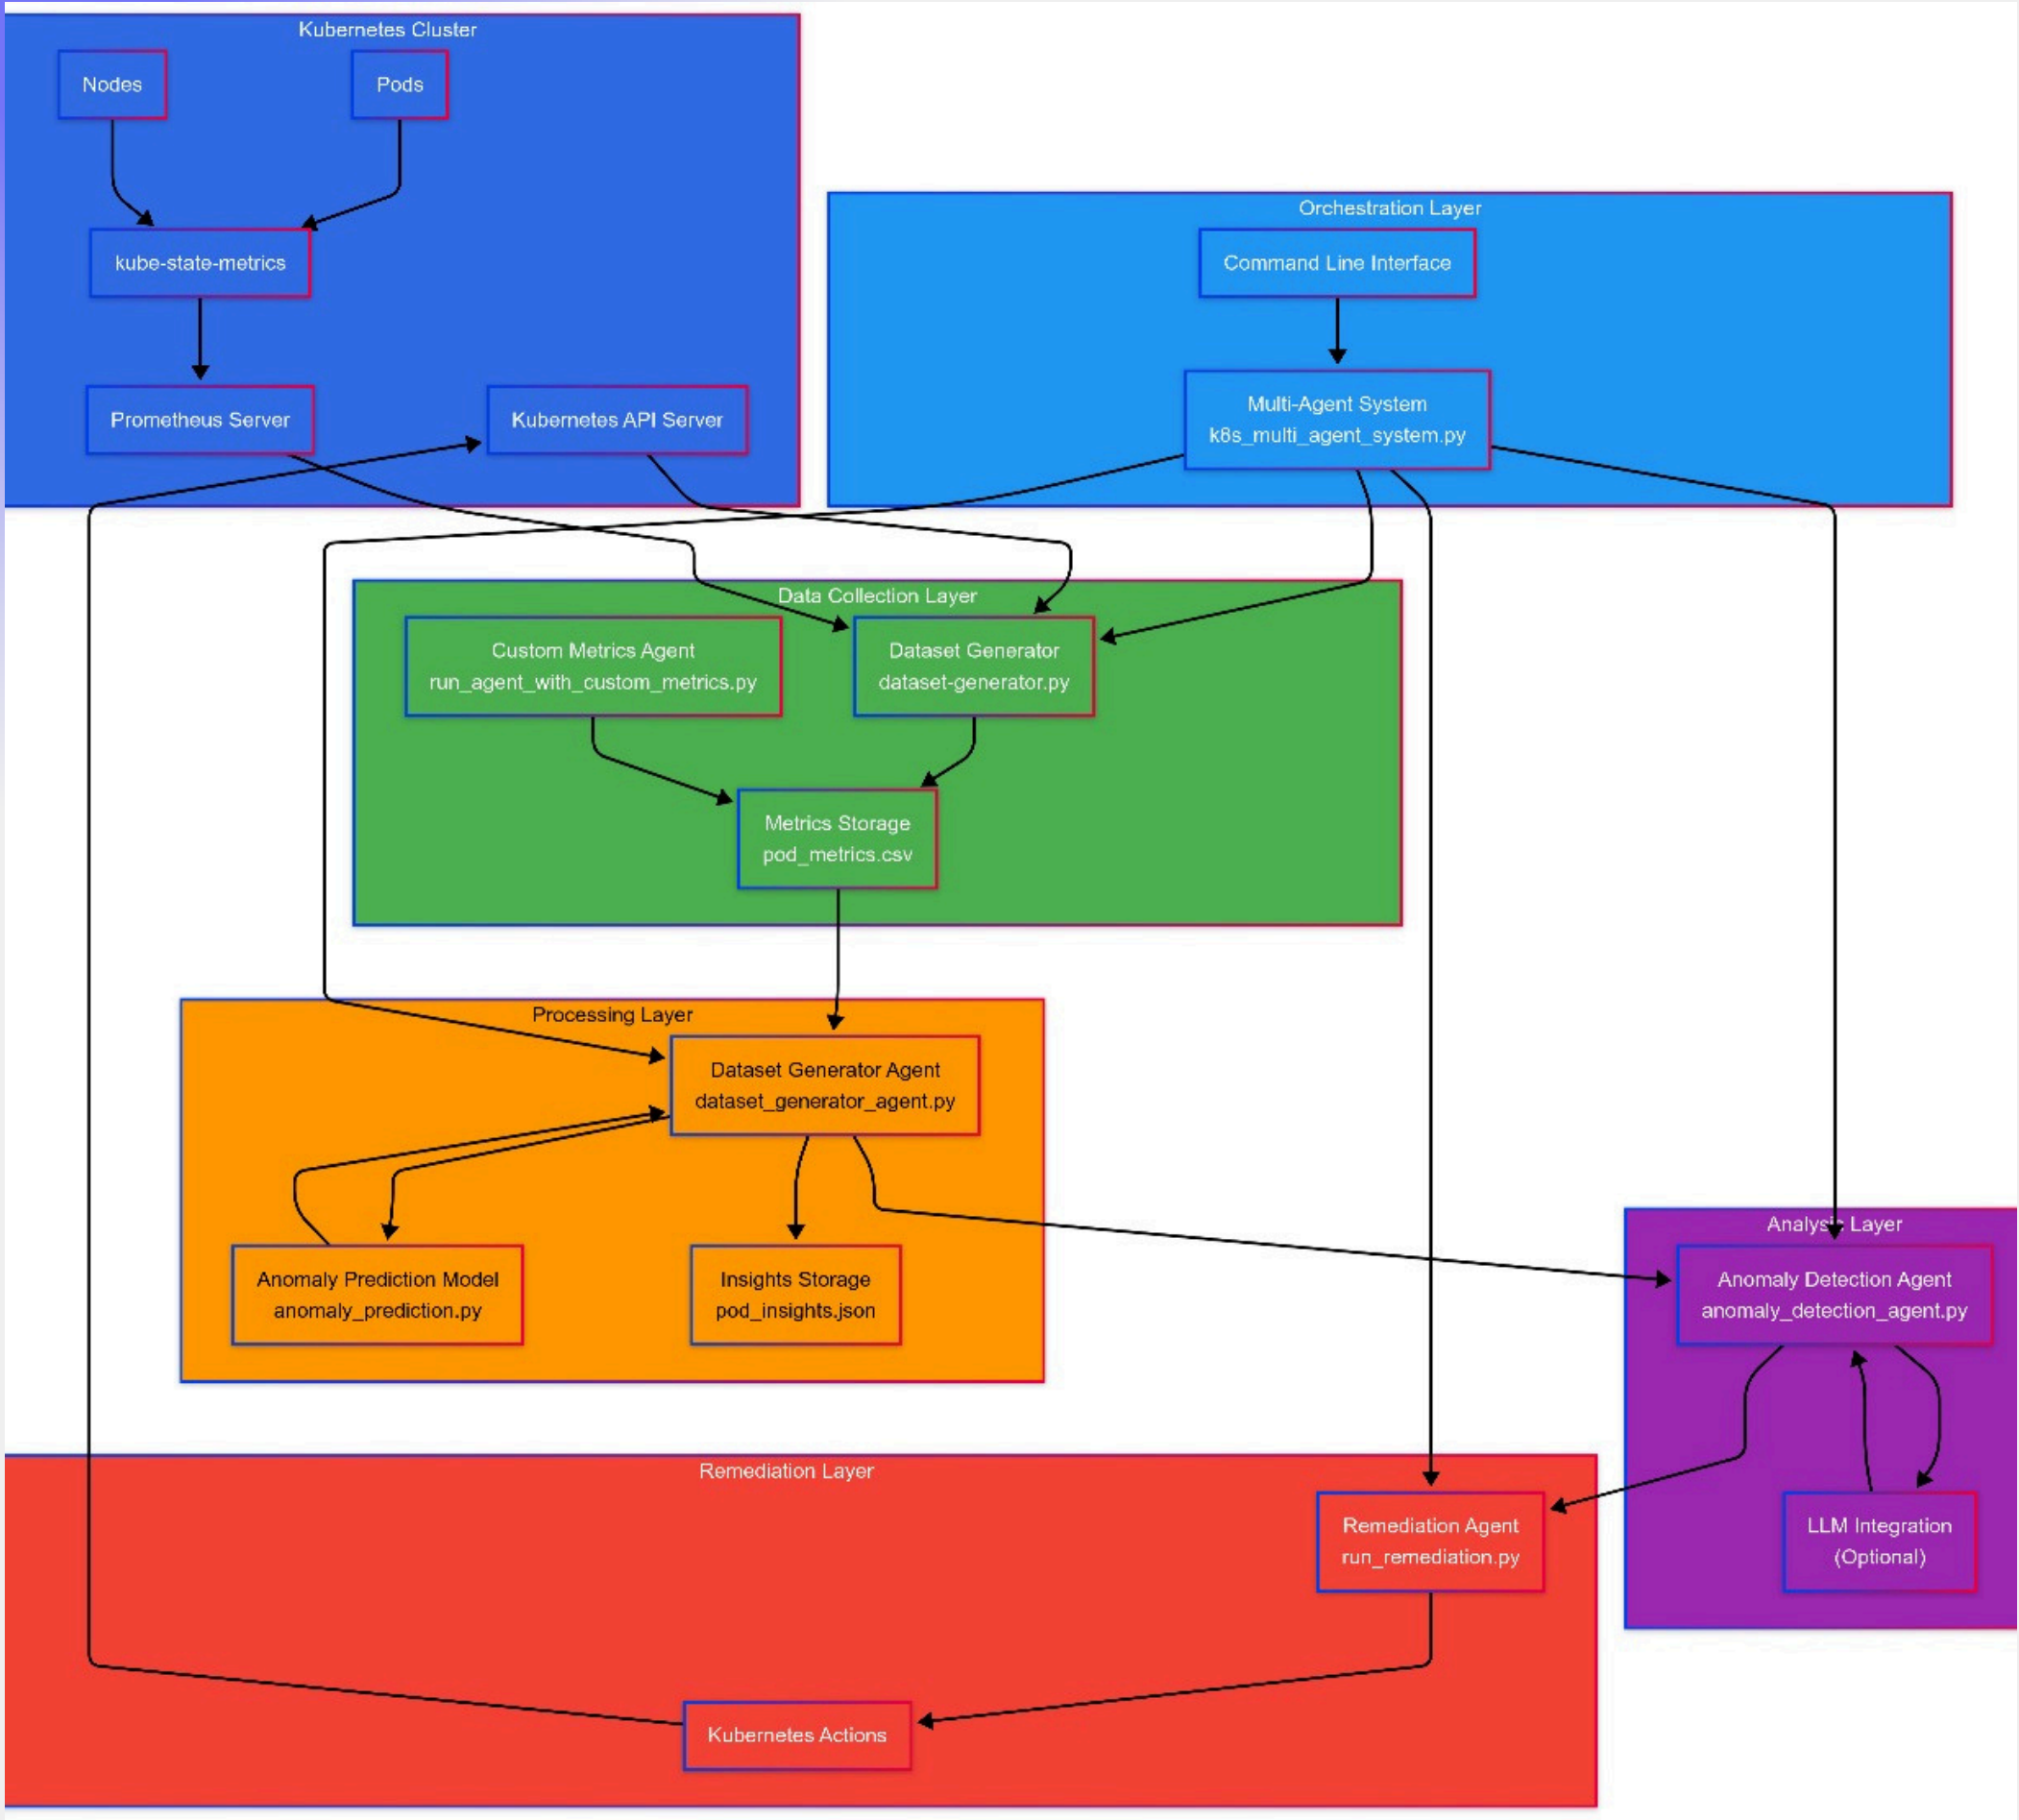
■ Multi - Agent System

- Coordinator-Diagnostic-Action agent architecture
- Auto-restart pods and scale resources
- KEDA integration for event-driven autoscaling
- 60% reduction in downtime
- Self-learning remediation strategy selection

■ Advanced Observability

- Integrated Prometheus and Grafana monitoring
- Real-time monitoring dashboards
- 40% faster incident resolution
- Single-pane visibility across cluster

Scope & Innovation



Architecture

05

Built With

Frameworks & Tools:

- Python, FastAPI, Prometheus, OpenTelemetry, Minikube, Docker, Kubernetes.

Monitoring Stack:

- Prometheus + Grafana for metrics, Node Exporter for system stats.

AI/ML Models:

- LSTM (for anomaly detection & prediction using metrics).

Remediation Layer:

- Custom agent using NVIDIA API to restart pods, auto scale, and manage resources.

Capital W

-

DevTrails
Hackathon 2025

06

07 Conclusion

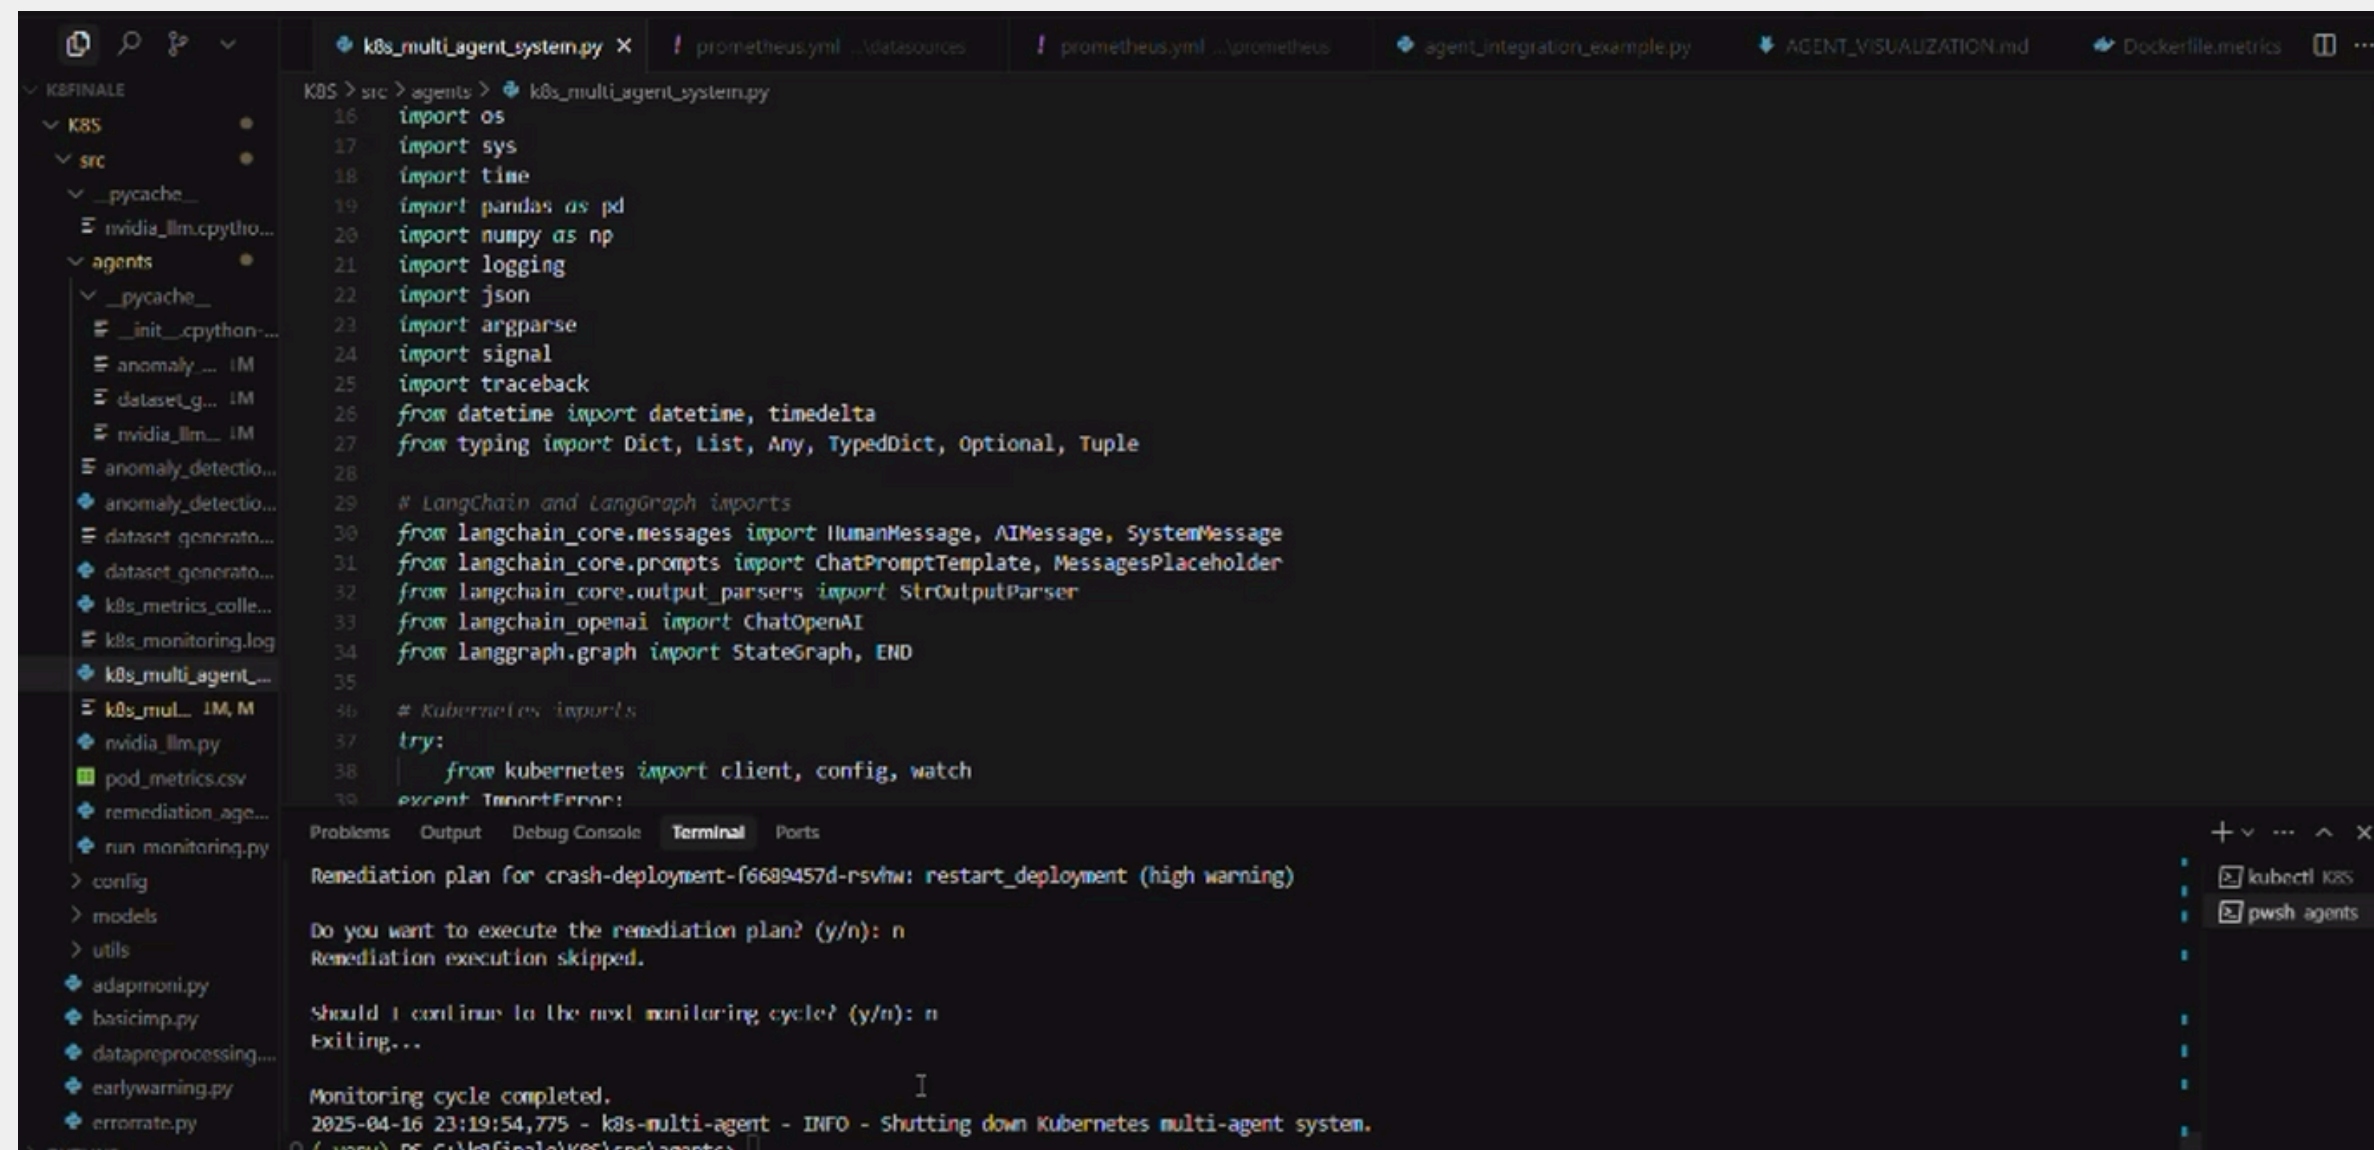
■ Key Benefits:

- Proactive failure detection in Kubernetes clusters.
- Automated remediation reduces manual DevOps effort.
- Scalable architecture using lightweight agents and Prometheus.
- Adaptable to various environments with minimal overhead.

■ Challenges & Mitigation:

- Limited hardware for real-time inference → Used lightweight ML models.
- Noisy or incomplete metrics → Injected synthetic anomalies to enrich training data.
- Frequent false alerts → Tuned model thresholds using historical cluster behavior.

Thank You



The screenshot displays a VS Code editor with a project named 'K8S' under the 'K8S' directory. The file explorer on the left shows a directory structure including 'src', 'agents', and various Python files. The main editor window shows the file 'k8s_multi_agent_system.py' with the following code:

```
K8S > src > agents > k8s_multi_agent_system.py
16 import os
17 import sys
18 import time
19 import pandas as pd
20 import numpy as np
21 import logging
22 import json
23 import argparse
24 import signal
25 import traceback
26 from datetime import datetime, timedelta
27 from typing import Dict, List, Any, TypedDict, Optional, Tuple
28
29 # LangChain and LangGraph imports
30 from langchain_core.messages import HumanMessage, AIMessage, SystemMessage
31 from langchain_core.prompts import ChatPromptTemplate, MessagesPlaceholder
32 from langchain_core.output_parsers import StrOutputParser
33 from langchain_openai import ChatOpenAI
34 from langgraph.graph import StateGraph, END
35
36 # Kubernetes imports
37 try:
38     from kubernetes import client, config, watch
39 except ImportError:
```

The terminal at the bottom shows the execution of the script, which includes a remediation plan for a crash-deployment, a prompt to execute the plan, and a completion message.

```
Remediation plan for crash-deployment-f6689457d-rsvhw: restart_deployment (high warning)
Do you want to execute the remediation plan? (y/n): n
Remediation execution skipped.
Should I continue to the next monitoring cycle? (y/n): n
Exiting...
Monitoring cycle completed.
2025-04-16 23:19:54,775 - k8s-multi-agent - INFO - Shutting down Kubernetes multi-agent system.
```

Video Explanation